



National and Kapodistrian  
University of Athens

# Shared Memory & Semaphores

Alex Delis

Univ. of Athens

November 2024

## IPCs (System V)

---

- Three types of IPCs:
  - Message Queues (not discussed)
  - Shared Memory
  - Semaphores (POSIX Interface)
- Each IPC structure is referred to by a non-negative integer identifier.
  - When an IPC is created, the program responsible for this creation provides a key of type `key_t`.
  - The Operating System converts this key into an IPC identifier.

# Keys in the IPC Client-Server Paradigm

---

⇒ Keys can be created in **three ways**:

1. The “server” program creates a new structure by specifying a private key that is `IPC_PRIVATE`.
  - Client has to **become explicitly aware** of this private key.
  - This is often accomplished with the help of a file generated by the server and then looked-up by the client.
2. Server and client **do agree** on a key value (often defined and hard-coded in the header).
3. Server and client can agree on a pathname to an existing file in the file system AND a project-ID (0..255) and then call `ftok()` to **convert** these two values into a **unique** key!

# Keys

---

- Keys help identify resources and offer access to the internal structures of the 3 IPC mechanisms (through systems calls):

```
struct msqid_ds    // for message queues
struct shmid_ds    // for shared segments
struct semid_ds    // for semaphores
```

- Wrongly accessing resources returns -1
- Access rights for IPC mechanisms: read/write stored in struct `ipc_perm`
- Included header files:

```
#include <sys/ipc.h>
#include <sys/types.h>
```

## The `ftok()` system call

---

- converts a pathname and a project identifier to a (System V) IPC-key

- `key_t ftok(const char *pathname, int proj_id)`


- Calling the `ftok()`:

```
if ( (thekey=ftok("/tmp/ad.tempfile", 23)) == -1)
    perror("Cannot create key from /tmp/ad.tempfile");
```

- The file `/tmp/ad.tempfile` must be accessible by the invoking process.

# Shared Memory

- A **shared memory region** is a portion of physical memory that is shared by multiple processes.



- In this region, structures can be set up by processes and others may read/write on them.
- Synchronization among processes using the segment (if required) is achieved with the help of **semaphores**.

## Creating a shared segment with `shmget()`

---

```
#include <sys/ipc.h>
#include <sys/shm.h>

int shmget(key_t key, size_t size, int shmflg)
```

- returns the **identifier** of the shared memory segment associated with the value of the argument `key`.
- the returned **size** of the segment is equal to `size` rounded up to a multiple of `PAGE_SIZE`.
- `shmflg` helps designate the access rights for the segment (`IPC_CREAT` and `IPC_EXCL` are used in a way similar to that of message queues).
- If `shmflg` specifies *both* `IPC_CREAT` and `IPC_EXCL` and a shared memory segment already exists for `key`, then `shmget()` fails with `errno` set to `EEXIST`.

## Attach- and Detach-ing a segment: `shmat()` / `shmdt()`

---

```
void *shmat(int shmid, const void *shmaddr, int shmflg)
```

- attaches the shared memory segment identified by `shmid` to the address space of the calling process.
- If `shmaddr` is `NULL`, the OS chooses a suitable (unused) address at which to attach the segment (frequent choice).
- Otherwise, `shmaddr` must be a page-aligned address at which the attach occurs.

```
int shmdt(const void *shmaddr)
```

- detaches the shared memory segment located at the address specified by `shmaddr` from the address space of the calling process.



# The system call `shmctl()`

---

```
int shmctl(int shmid, int cmd, struct shmid_ds *buf)
```

- performs the control operation specified by `cmd` on the shared memory segment whose identifier is given in `shmid`.
- The `buf` argument is a pointer to a `shmid_ds` structure:

```
struct shmid_ds {  
    struct ipc_perm shm_perm;    /* Ownership and permissions */  
    size_t          shm_segsz;   /* Size of segment (bytes) */  
    time_t          shm_atime;   /* Last attach time */  
    time_t          shm_dtime;   /* Last detach time */  
    time_t          shm_ctime;   /* Last change time */  
    pid_t           shm_cpid;    /* PID of creator */  
    pid_t           shm_lpid;    /* PID of last shmat(2)/shmdt(2) */  
    shmatt_t        shm_nattch;  /* No. of current attaches */  
    ...  
};
```

## The system call `shmctl()`

---

Usual values for `cmd` are:

- `IPC_STAT`: copy information from the kernel data structure associated with `shmid` into the `shmid_ds` structure pointed to by `buf`.
- `IPC_SET`: write the value of some member of the `shmid_ds` structure pointed to by `buf` to the kernel data structure associated with this shared memory segment, updating also its `shm_ctime` member.
- `IPC_RMID`: mark the segment to be destroyed. The segment will be destroyed after the last process detaches it (i.e., `shm_nattch` is zero).

# Use Cases of Calls

---

- Only one process creates the segment:

```
int id;  
id = shmget(IPC_PRIVATE, 10, 0666);  
if ( id == -1 ) perror("Creating");
```

- Every (interested) process attaches the segment:

```
int *mem;  
mem = (int *) shmat (id, (void *)0, 0);  
if ( (int)mem == -1 ) perror("Attachment");
```

- Every process detaches the segment:

```
int err;  
err = shmdt((void *)mem);  
if ( err == -1 ) perror("Detachment");
```

- Only one process has to remove the segment:

```
int err;  
err = shmctl(id, IPC_RMID, 0);  
if ( err == -1 ) perror("Removal");
```

## Creating and accessing shared memory (shareMem1.c)

---

```
#include <stdio.h>
#include <sys/types.h>
#include <sys/ipc.h>
#include <sys/shm.h>

int main(int argc, char **argv){
    int id=0, err=0;
    int *mem;

    id = shmget(IPC_PRIVATE,10,0666); /* Make shared memory segment */
    if (id == -1) perror ("Creation");
    else printf("Allocated. \u00d\n", (int)id);

    mem = (int *) shmat(id, (void*)0, 0); /* Attach the segment */
    if (*(int *) mem == -1) perror("Attachment.");
    else printf("Attached. \u00Mem\u00contents \u00d\n", *mem);

    *mem=1; /* Give it initial value */
    printf("Start \u00other \u00process. \u00>"); getchar();

    printf("mem \u00is \u00now \u00d\n", *mem); /* Print out new value */

    err = shmctl(id, IPC_RMID, 0); /* Remove segment */
    if (err == -1) perror ("Removal.");
    else printf("Removed. \u00d\n", (int)(err));
    return 0;
}
```

## Creating and accessing shared memory (shareMem2.c)

---

```
#include <stdio.h>
#include <stdlib.h>
#include <sys/types.h>
#include <sys/ipc.h>
#include <sys/shm.h>

int main(int argc, char **argv) {
    int id, err;
    int *mem;

    if (argc <= 1) { printf("Need shared memory id.\n"); exit(1); }

    sscanf(argv[1], "%d", &id); /* Get id from command line. */
    printf("Id is %d\n", id);

    mem = (int *) shmat(id, (void*) 0, 0); /* Attach the segment */
    if ((int) mem == -1) perror("Attachment.");
    else printf("Attached. Mem contents %d\n", *mem);

    *mem=2; /* Give it a different value */
    printf("Changed mem is now %d\n", *mem);

    err = shmdt((void *) mem); /* Detach segment */
    if (err == -1) perror ("Detachment.");
    else printf("Detachment %d\n", err);
    return 0;
}
```

## Running the two programs:

---

- Starting off with executing "shareMem1":

```
ad@haiku:~/SharedSegments$ ./shareMem1
Allocated. 1769489
Attached. Mem contents 0
Start other process. >
```

- Executing "shareMem2":

```
ad@haiku:~/SharedSegments$ ./shareMem2 1769489
Id is 1769489
Attached. Mem contents 1
Changed mem is now 2
Detachment 0
ad@haiku:~/SharedSegments$
```

- Providing the final input to "shareMem1":

```
Start other process. >s
mem is now 2
Removed. 0
ad@haiku:~/SharedSegments$
```

# Semaphores

---

- **Fundamental mechanism** that facilitates **synchronization and coordinated accessing** of resources placed in shared memory.
- A semaphore is an integer whose value is **never allowed** to fall below zero.
- *Two operations* can be atomically performed on a semaphore:
  - **increment** the semaphore value by one (UP or V() ala Dijkstra).
  - **decrement** a semaphore value by one (DOWN or P() ala Dijkstra).  
If the value of semaphore is currently zero, then the invoking process will block until the value becomes greater than zero.

# POSIX Semaphores

---

```
#include <semaphore.h>
```

- `sem_init`, `sem_destroy`, `sem_post`, `sem_wait`, `sem_trywait`

```
int sem_init(sem_t *sem, int pshared, unsigned int value);
```

- The above initializes a semaphore.
- Compile either with `-lrt` or `-lpthread`
- `pshared` indicates whether this semaphore is to be shared between the threads of a process, or between processes:
  - **zero**: semaphore is shared between the **threads of a process**; should be located at an address visible to **all threads**.
  - **non-zero**: semaphore is shared **among processes**.



## POSIX Semaphore Operations

---

- `sem_wait()`, `sem_trywait()`
  - `int sem_wait(sem_t *sem);`  
`int sem_trywait(sem_t *sem);`
  - Perform P(s) operation.
  - `sem_wait` blocks; `sem_trywait` will fail rather than block.
- `sem_post()`
  - `int sem_post(sem_t *sem)`
  - Performs V(s) operation.
- `sem_destroy()`
  - `int sem_destroy(sem_t *sem);`
  - Destroys a semaphore.

## Creating and using a POSIX Semaphore

---

```
#include <stdio.h>
#include <stdlib.h>
#include <semaphore.h>
#include <sys/types.h>
#include <sys/ipc.h>

extern int errno;

int main(int argc, char **argv)
{
    sem_t sp; int retval;

    /* Initialize the semaphore. */
    retval = sem_init(&sp,1,2);
    if (retval != 0) {
        perror("Couldn't initialize."); exit(3); }

    retval = sem_trywait(&sp);
    printf("Did trywait. Returned %d>\n",retval); getchar();

    retval = sem_trywait(&sp);
    printf("Did trywait. Returned %d>\n",retval); getchar();

    retval = sem_trywait(&sp);
    printf("Did trywait. Returned %d>\n",retval); getchar();

    sem_destroy(&sp);
    return 0;
}
```

## Executing the Program

---

```
ad@ad-desktop:~/src/PosixSems$ ./semtest
Did trywait. Returned 0 >

Did trywait. Returned 0 >

Did trywait. Returned -1 >

ad@ad-desktop:~/src/PosixSems$
```

## Example of Shared Memory & Semaphore: semtest3.c

---

```
/* semtest3.c: POSIX Semaphore test example using shared memory */
#include <stdio.h>
#include <stdlib.h>
#include <semaphore.h>
#include <sys/types.h>
#include <sys/ipc.h>
#include <sys/shm.h>

#define SEGMENTSIZE sizeof(sem_t)
#define SEGMENTPERM 0666

int main(int argc, char **argv)
{
    sem_t *sp;
    int retval;
    int id, err;

    /* Make shared memory segment. */
    id = shmget(IPC_PRIVATE, SEGMENTSIZE, SEGMENTPERM);
    if (id == (void *) -1) perror("Creation");
    else printf("Allocated_%d\n", id);

    /* Attach the segment. */
    sp = (sem_t *) shmat(id, (void*) 0, 0);
    if ( sp == (void *) -1) { perror("Attachment."); exit(2);}
```

## Example: semtest3.c

---

```
/* Initialize the semaphore. */
retval = sem_init(&sp,1,2);
if (retval != 0) {
    perror("Couldn't initialize.");
    exit(3);
}
retval = sem_trywait(&sp);
printf("Did trywait. Returned %d>\n", retval); getchar();

retval = sem_trywait(&sp);
printf("Did trywait. Returned %d>\n", retval); getchar();

retval = sem_trywait(&sp);
printf("Did trywait. Returned %d>\n", retval); getchar();

sem_destroy(&sp);

/* Remove segment. */
err = shmctl(id, IPC_RMID, 0);
if (err == -1) perror("Removal.");
else printf("Removed. %d\n",err);
return 0;
```

```
}
```

## Example: semtest3a.c

---

```
/* semtest3a.c POSIX Semaphore test example using shared memory */
#include <stdio.h>
#include <stdlib.h>
#include <semaphore.h>
#include <sys/types.h>
#include <sys/ipc.h>
#include <sys/shm.h>
#define SEGMENTSZIE sizeof(sem_t)
#define SEGMENTPERM 0666

extern int errno;

int main(int argc, char **argv)
{
    sem_t *sp; int retval; int id, err;

    if (argc <= 1) { printf("Need shm id.\n"); exit(1); }

    /* Get id from command line. */
    sscanf(argv[1], "%d", &id);
    printf("Allocated %d\n", id);

    /* Attach the segment. */
    sp = (sem_t *) shmat(id, (void *) 0, 0);
    if (sp == (void *) -1) { perror("Attachment."); exit(2); }

    /* Initialize the semaphore. */
    retval = sem_init(sp, 1, 1);
    if (retval != 0) { perror("Couldn't initialize."); exit(3); }
```

## Example: semtest3.c

---

```
retval = sem_trywait(sp);
printf("Did_trywait.Returned_%d>\n", retval); getchar();

retval = sem_trywait(sp);
printf("Did_trywait.Returned_%d>\n", retval); getchar();

retval = sem_trywait(sp);
printf("Did_trywait.Returned_%d>\n", retval); getchar();

/* Remove segment. */
err = shmdt((void *) sp);
if (err == -1) perror ("Detachment.");
return 0;
```

```
}
```

## Running the two programs: semtest and semtest3a

---

- Starting off with executing "semtest":

```
ad@serifos:~/Recitation3/src$ ./semtest3
Allocated 14549024
Did trywait. Returned 0 >
```

- Executing "semtest3a" in another tty:

```
ad@serifos:~/Recitation3/src$ ./semtest3a 14549024
Allocated 14549024
Did trywait. Returned 0 >

Did trywait. Returned -1 >
```

- Continue with "semtest3":

```
Did trywait. Returned -1 >

Did trywait. Returned -1 >
```



## Running the two programs: semtest and semtest3a

---

- Continue with "semtest3a":

```
Did trywait. Returned -1 >
```

- Follow up with "semtest3":

```
Removed. 0  
ad@serifos:~/Recitation3/src$
```

- Finish with "semtest3a":

```
ad@serifos:~/Recitation3/src$
```

## Initialize and Open a **named Semaphore**

---

```
sem_t *sem_open(const char *name, int oflag);  
sem_t *sem_open(const char *name, int oflag,  
                mode_t mode, unsigned int value);
```

- creates a new POSIX semaphore OR opens an existing semaphore whose name is `name`.
- `oflag` specifies flags that control the operation of the call
  - `O_CREAT` creates the semaphore;
  - provided that both `O_CREAT` and `O_EXCL` are specified, an error is returned if a semaphore with `name` already exists.
- if `oflag` is `O_CREAT` then **2 more arguments** have to be used:
  - `mode` specifies the permissions to be placed on the new semaphore.
  - `value` specifies the initial value for the new semaphore.

## More on Named POSIX Semaphores

---

- A named semaphore is identified by a (persistent) name that has the form `/this_is_a_sample_named_semaphore`.
  - consists of an initial slash followed by a (large) number of character (but no slashes).
- If you want to “see”(list) all **named sempahores** in your (Linux) system look at directory `/dev/shm`

## More on Named POSIX Semaphores

---

```
int sem_close(sem_t *sem)
```

- closes the named semaphore referred to by *sem* freeing the system resources the invoking process has used.

```
int sem_unlink(const char *name)
```

- removes the named semaphore in question.

```
int sem_getvalue(sem_t *sem, int *sval)
```

- obtains the current value of semaphore..
- the **cheater** API-call!

# Named POSIX Semaphore

---

```
#include      <stdio.h>
...
#include      <sys/stat.h>
#include      <semaphore.h>

int main(int argc, char *argv[]){
    const char *semname;
    int op=0; int val=0;

    if (argc==3) {
        semname=argv[1]; op=atoi(argv[2]);
    }
    else {
        printf("usage: _nameSem _nameOfSem _Operation\n"); exit(1);
    }

    sem_t *sem=sem_open(semname, O_CREAT|O_EXCL, S_IRUSR|S_IWUSR, 0);

    if (sem!= SEM_FAILED)
        printf("created _new _semaphore!\n");
    else if (errno== EEXIST ) {
        printf("semaphore _appears _to _exist _already!\n");
        sem = sem_open(semname, 0);
    }
    else ;

    assert(sem != SEM_FAILED);
    sem_getvalue(sem, &val);
    printf("semaphore's _before _action _value _is _%d\n", val);
```

## Named Posix Semaphore

---

```
if ( op == 1 ) {
    printf("incrementing semaphore\n");
    sem_post(sem);
}
else if ( op == -1 ) {
    printf("decrementing semaphore\n");
    sem_wait(sem);
}
else if ( op == 2 ){
    printf("clearing up named semaphore\n");
    sem_close(sem); // close the sem
    sem_unlink(semname); // remove it from system
    exit(1);
}
else    printf("not defined operation!\n");

sem_getvalue(sem, &val);
printf("semaphore's current value is %d\n", val);
sem_close(sem);
return(0);
}
```

# Execution Outcome

```
ad@serifos:~/PosixSems$ ls /dev/shm/
pulse-shm-1024070233  pulse-shm-1294442337  pulse-shm-2927836935
pulse-shm-1274848112  pulse-shm-2305588894  pulse-shm-3888866544
ad@serifos:~/PosixSems$ ./namedSem /delis 1
created new semaphore!
semaphore's before action value is 0
incrementing semaphore
semaphore's current value is 1
ad@serifos:~/PosixSems$ ls /dev/shm/
pulse-shm-1024070233  pulse-shm-1294442337  pulse-shm-2927836935  sem.delis
pulse-shm-1274848112  pulse-shm-2305588894  pulse-shm-3888866544
ad@serifos:~/PosixSems$ ./namedSem /delis -1
semaphore appears to exist already!
semaphore's before action value is 1
decrementing semaphore
semaphore's current value is 0
ad@serifos:~/PosixSems$ ./namedSem /delis 2
semaphore appears to exist already!
semaphore's before action value is 0
clearing up named semaphore
ad@serifos:~/PosixSems$ ls /dev/shm/
pulse-shm-1024070233  pulse-shm-1294442337  pulse-shm-2927836935
pulse-shm-1274848112  pulse-shm-2305588894  pulse-shm-3888866544
ad@serifos:~/PosixSems$ ./namedSem /delis 1
created new semaphore!
semaphore's before action value is 0
incrementing semaphore
semaphore's current value is 1
```

## Execution Outcome

---

```
ad@serifos:~/PosixSems$ ./namedSem /delis 1
semaphore appears to exist already!
semaphore's before_action_value is 1
incrementing semaphore
semaphore's current value is 2
ad@serifos:~/PosixSems$ ls /dev/shm/
pulse-shm-1024070233  pulse-shm-1294442337  pulse-shm-2927836935  sem.delis
pulse-shm-1274848112  pulse-shm-2305588894  pulse-shm-3888866544
ad@serifos:~/PosixSems$ ./namedSem /delis -1
semaphore appears to exist already!
semaphore's before_action_value is 2
decrementing semaphore
semaphore's current value is 1
ad@serifos:~/PosixSems$ ./namedSem /delis -1
semaphore appears to exist already!
semaphore's before_action_value is 1
decrementing semaphore
semaphore's current value is 0
ad@serifos:~/PosixSems$ ./namedSem /delis 2
semaphore appears to exist already!
semaphore's before_action_value is 0
clearing up named semaphore
ad@serifos:~/PosixSems$ ls /dev/shm/
pulse-shm-1024070233  pulse-shm-1294442337  pulse-shm-2927836935
pulse-shm-1274848112  pulse-shm-2305588894  pulse-shm-3888866544
ad@serifos:~/PosixSems$
```